

08 — Sprint Tracker (global numbering)

Owner: `fce-lead-systems-architect`. Repo home: `docs/handoff/08_sprint-tracker.md`. This is the single reference for: which sprint we are in, what it produces, where it runs, and how long it should take. Update the Status column at the end of every sprint chat — this file is the pointer, the block docs (`14`, `15`) remain the detailed source of truth.

Numbering rule

Sprints are numbered globally and consecutively: each mission block owns two. M1 = Sprints 1-2, M2 = Sprints 3-4, M3 = Sprints 5-6, M4 = Sprints 7-8, M5 = Sprints 9-10, M6 = Sprints 11-12, M7 = Sprints 13-14 (TRL 1-3 scope), M8 = Sprints 15-16, M9 = Sprints 17-18 (design-only at TRL 1-3). Older documents that say "M2 Sprint 2" mean global Sprint 4 — the mapping is always: $global = (block\ number - 1) \times 2 + block\text{-}local\ sprint$.

Effort estimates

All durations are ENGINEERING JUDGMENT planning estimates for one developer working with Claude (chat for design/review, Claude Code for repo/code work), in focused working days. They are internal planning figures — never schedule commitments, never proposal claims. Rule of thumb: design/spec sprints run 0.5-1 day; PoC/code sprints run 1-2 days; the fixture sprint runs 1-2 days including data download and trimming.

Tracker

Sprint	Block	Objective	Key outputs	Runs in	Est. effort	G
1	M1	Replace paraphrased outcome anchors with verbatim solicitation text; draft finalized RTM rows with acceptance criteria	Verbatim outcome registry in <code>02</code> / <code>03</code> with Canada.ca citation; 17 RTM rows (16 original + FCE-REQ-KRN-011 added by accountability review); live-source verification evidence	Chat (this Project) + Claude Code apply	0.5-1 day	—

Sprint	Block	Objective	Key outputs	Runs in	Est. effort	G
2	M1	Coverage audit: ESS 6/6, DES 4/4 (FCE-DES-01/03 confirmed), all 17 rows singular/testable/correct shall-should, verification methods reviewed	Coverage report (EVD-M1), gap list, GATE-A recommendation	Chat, then Claude Code commit	0.5-1 day	GA
3	M2	Ratify FCE-DR-SCH-001/002 and freeze the 15-field schema at v0.2.0; confirm taxonomy mapping approach (OPEN-02); fail-closed at G2	Schema v1 freeze record, provenance model spec, validation-rule list, PoC file plan	Chat	0.5-1 day	—
4	M2	Implement minimal schema-validation PoC and tests (valid object, missing field, malformed, source-supplied policy_binding_state)	Schema validator PoC, fixtures, test output (EVD-M2)	Claude Code	1-2 days	G (P)
5	M3	Finalize deterministic default-deny PDP/PEP/PAP/PIP model: 11 actions, reason codes incl. RC-008, deny-overrides, B1 PIP auth, B2 override envelope	Policy model v1, rule examples, evaluator file plan	Chat	0.5-1 day	—
6	M3	Implement machine-readable policy fixture + local deterministic evaluator returning permit/restrict/block/quarantine/review	Policy evaluator PoC, determinism tests, PIP-spoof and override red-team tests (EVD-M3)	Claude Code	1-2 days	G (P)
7	M4	Finalize 18-field audit schema, 9 event classes, hash-chain semantics, replay spec, export/manifest shape	Audit schema v1, replay spec, PoC file plan	Chat	0.5-1 day	—
8	M4	Implement JSONL audit emission + provenance parent-link capture for accepted/rejected/transformed/fused objects	Audit writer PoC, sample JSONL, replay check (EVD-M4)	Claude Code	1-2 days	G cl w S)

Sprint	Block	Objective	Key outputs	Runs in	Est. effort	G
						4, al
9	M5	Finalize fusion-kernel interface, high-water-mark propagation, explicit-permit merge check, segregation-on-block — tested against FCE-REQ-KRN-011	Fusion-kernel spec, merge decision model	Chat	0.5-1 day	—
10	M5	Implement merge evaluation over 2+ parents: permitted same-domain merge + blocked cross-domain merge with RC-003 and audit	No-merge PoC, blocked-merge evidence (EVD-M5)	Claude Code	1-2 days	G D (f
11	M6	Finalize four scenario specs (Joint ISR, Maritime, Tactical Edge, UAV) with embedded conflicts; confirm source choices under OPEN-04	Scenario library, trim protocol, split plan, expected dispositions	Chat	0.5-1 day	—
12	M6	Download, trim, manifest, split, and seal fixtures: 2 approved open-source families + synthetic red-team variants (tampered, malformed, stale, PIP spoof, pre-marking, unauthorized merge)	Source manifest, trim report, calibration + sealed held-out fixture sets, coverage matrix (EVD-M6)	Claude Code	1-2 days	G C
13	M7	Finalize bidirectional V&V matrix across unit/integration/property-based/explainability/red-team classes; every requirement ≥ 1 test; every GDR guard has a rejection test	V&V matrix, test-harness file plan (target file <u>docs/17_vv-plan.md</u>)	Chat	0.5-1 day	—

Sprint	Block	Objective	Key outputs	Runs in	Est. effort	G
14	M7	Implement and run the two-layer harness: code-correctness tests, then the sealed held-out validation run; report negative results verbatim	Runnable harness, guard rejection tests, code-correctness report, held-out validation report, coverage report (EVD-M7)	Claude Code	1-2 days	G D w Sj 10 d
15	M8	Define benchmark metric set (per-gate + end-to-end latency, throughput, memory, CPU/GPU, thermal, power) — all TARGET; six degraded-mode classes with fail-closed responses	Benchmark plan (TARGET-labelled)	Chat	0.5-1 day	—
16	M8	Baseline-vs-candidate protocol on named hardware; NVIDIA decision records DR-001...007 outside the compliance path	Degraded-mode protocol, NVIDIA decision records	Chat (execution deferred to TRL 4-5 approval)	0.5-1 day	G
17	M9	Evidence index (EVD-*) per TRL band; TRL exit-criteria gap audit	Evidence pack index, gap report	Chat	0.5-1 day	—
18	M9	Compliance matrix (clause → evidence), claim screen, red-team claim audit, website-safe and proposal-safe language	Compliance matrix, proposal-ready sections	Chat	1 day	G

Roll-up

TRL 1-3 execution scope is Sprints 1-14 (M1-M7): roughly 11-18 focused days at the estimates above — call it 3-4 calendar weeks for one developer at sustained part-time pace, before M8/M9 design sprints. ENGINEERING JUDGMENT; re-estimate after Sprint 4, the first code sprint, which calibrates the PoC velocity.

Standing rules per sprint chat

One block per chat; one sprint per output; never start the next global sprint number in the same chat. Every sprint chat opens by reading this tracker plus the block entry in (15), runs

its role blocks per the one-chat rule, ends with the role handoff plus a one-line tracker status update to commit via Claude Code. Chat drafts REPO-UPDATE notes; Claude Code writes and runs. A sprint is DONE only when its evidence artifact (EVD-*) exists at the referenced path and this tracker's Status column is updated in the repo.

Facts / Assumptions / Judgment / Uncertainty

- **Facts:** block objectives, outputs, gates, and dependencies are drawn from `docs/14` and `docs/15` at commit 4637107; Sprint 1 completion state is verified against the repo.
- **Assumptions:** OPEN-04 (data sources) and OPEN-02 (taxonomy mapping) resolve without schema rework; one developer plus Claude at part-time pace.
- **Judgment:** the global numbering convention, all effort estimates, and the suggested USGS + Sentinel-2 source pairing.
- **Uncertainty:** PoC sprint velocity unknown until Sprint 4; OPEN-03 hardware timing affects only Sprints 15-16.